

Transform iptables into a TCP load balancer

Task requirement

To create two new Nginx LB or two nginx and after creating, it needs to forward the 50% of Load to another LB.

iptables is a user-space utility program that allows a system administrator to configure the IP packet filter rules of the Linux kernel firewall.

Environment details

IPtable rules will be applicable for all types of OS.

UbuntuNAME="Ubuntu"
VERSION="20.04.4"

List of tools and technologies

- Nginx:- NGINX_VERSION=1.25.1

I have configure two vm and install nginx both vm

VM1 >>>> 192.168.122.123
sudo apt install nginx

VM2 >>>> 192.168.122.44
sudo apt install nginx

Vm1 >>>>Change the index.html

jai@jai-Standard-PC-Q35-ICH9-2009:~\$ cd /var/www/html/

jai@jai-Standard-PC-Q35-ICH9-2009:/var/www/html\$ ls
index.html

jai@jai-Standard-PC-Q35-ICH9-2009:/var/www/html\$ vim index.html

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Hello Deepak!</title>
</head>
<body>
  <h1>Hello Deepak!</h1>
</body>
```

</html>

Save file

:wq!

Vm2 >>>> Change the index.html

jai1@jai1-Standard-PC-Q35-ICH9-2009:/var/www/html\$ ls

index.nginx-debian.html

jai1@jai1-Standard-PC-Q35-ICH9-2009:/var/www/html\$ sudo vim index.nginx-debian.html

```
<!DOCTYPE html>
```

```
<html lang="en">
```

```
<head>
```

```
    <meta charset="UTF-8">
```

```
    <meta name="viewport" content="width=device-width, initial-scale=1.0">
```

```
    <title>Hello jaiprakash!</title>
```

```
</head>
```

```
<body>
```

```
    <h1>Hello jaiprakash!</h1>
```

```
</body>
```

```
</html>
```

>>>> sudo apt install curl

I am able to curl Page from each other.

jai@jai-Standard-PC-Q35-ICH9-2009:~\$ curl http://192.168.122.123

```
<!DOCTYPE html>
```

```
<html lang="en">
```

```
<head>
```

```
    <meta charset="UTF-8">
```

```
    <meta name="viewport" content="width=device-width, initial-scale=1.0">
```

```
    <title>Hello Deepak!</title>
```

```
</head>
```

```
<body>
```

```
    <h1>Hello Deepak!</h1>
```

```
</body>
```

```
</html>
```

jai1@jai1-Standard-PC-Q35-ICH9-2009:/etc/nginx\$ curl http://192.168.122.44

```
<!DOCTYPE html>
```

```
<html lang="en">
```

```
<head>
```

```
    <meta charset="UTF-8">
```

```
    <meta name="viewport" content="width=device-width, initial-scale=1.0">
```

```
    <title>Hello jaiprakash!</title>
```

```
</head>
<body>
    <h1>Hello jaiprakash!</h1>
</body>
</html>
```

I am able to curl both from the Base Machine.

```
indianrenters@jai:~$ curl http://192.168.122.123
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Hello Deepak!</title>
</head>
<body>
  <h1>Hello Deepak!</h1>
</body>
</html>
```

```
indianrenters@jai:~$ curl http://192.168.122.44
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Hello jaiprakash!</title>
</head>
<body>
  <h1>Hello jaiprakash!</h1>
</body>
</html>
```

```
=====now work in Vm1 192.168.122.123=====
```

a) Enable IP Forwarding:

Activate IP forwarding to facilitate traffic routing between interfaces:

[illegible]

uncomment in the file below line.

```
net.ipv4.ip_forward=1
```

```
net.ipv4.conf.all.accept_redirects = 0
net.ipv4.conf.all.accept_source_route = 0
```

```
net.ipv4.conf.all.log_martians = 1
```

```
jai@jai-Standard-PC-Q35-ICH9-2009:~$ sudo sysctl -p
```

```
net.ipv4.ip_forward = 1
net.ipv4.conf.all.accept_redirects = 0
net.ipv4.conf.all.secure_redirects = 1
net.ipv4.conf.all.accept_source_route = 0
net.ipv4.conf.all.log_martians = 1
```

Modify source IP for forwarding: this command run vm1 source

```
jai@jai-Standard-PC-Q35-ICH9-2009:~$ sudo iptables -A POSTROUTING -t nat -p tcp -d
192.168.122.44 --dport 80 -j SNAT --to-source 192.168.122.123
```

Set default connection drop vm1

This implies that by default, any incoming traffic that doesn't match specific predefined rules will be dropped or rejected. In other words, if the firewall doesn't have a rule that explicitly allows the traffic to pass through, it will block it as a security measure. This approach follows the principle of allowing only known and authorised traffic, enhancing the network's security.

```
sudo iptables -t filter -P FORWARD DROP
```

Accept Traffic to the Server:

This involves creating rules that explicitly allow certain types of traffic to reach a designated server. For instance, if you have a server at IP address 192.168.122.44 and it's listening on port 80, you can configure the firewall to accept incoming traffic destined for that server's IP and port. This is done using firewall rules that specify the source, destination, protocol, and port of the allowed traffic.

Allow specific traffic from particular sources and to specific destinations (servers) on specific ports (--dport for destination port and --sport for source port). This creates a controlled and secure network environment where only the specified traffic is permitted to traverse the firewall.

This command on run vm1

```
jai@jai-Standard-PC-Q35-ICH9-2009:~$ sudo iptables -t filter -A FORWARD -d
192.168.122.44 -p tcp --dport 80 -j ACCEPT
jai@jai-Standard-PC-Q35-ICH9-2009:~$ sudo iptables -t filter -A FORWARD -s
192.168.122.44 -p tcp --sport 80 -j ACCEPT
```

This command run on vm2

```
jai1@jai1-Standard-PC-Q35-ICH9-2009:~$ sudo iptables -A INPUT -p tcp --dport 80 -j
ACCEPT
```

Apply Random Load Balancing and then Round Robin:-

A technique used in network environments to distribute incoming traffic across multiple servers in a random manner. This method is employed to optimize resource utilization, prevent overload on individual servers, and enhance the overall performance and reliability of a system.

In this case, incoming traffic destined for the IP address 192.168.122.123 on port 80 is subject to random distribution between one destinations ([192.168.122.44:80](#)). The --mode random option ensures that the traffic is distributed based on a randomized algorithm, with different probabilities assigned to each destination (--probability 0.33 and --probability 0.5 in this case).

The goal of applying random load balancing is to distribute traffic unpredictably, achieving a fair distribution of incoming requests among the specified destinations. This way, the servers can collectively handle the load more efficiently, minimizing the risk of overloading any single server and contributing to improved performance and fault tolerance.

This command run Vm1 (source)

```
sudo iptables -A PREROUTING -t nat -p tcp -d 192.168.122.123 --dport 80 \
-m statistic --mode random --probability 0.33 \
-j DNAT --to-destination 192.168.122.44:80
```

Capture Specific Port Traffic:

Use tcpdump to monitor network traffic on a specific port:-

```
tcpdump -i enp1s0 port 80 -n
```

Output of TCP Dump

1st curl hit

```
11:01:05.899106 IP 192.168.122.1.58216 > 192.168.122.123.80: Flags [S], seq
2182467641, win 64240, options [mss 1460,sackOK,TS val 2517725285 ecr 0,nop,wscale
7], length 0
11:01:05.899839 IP 192.168.122.123.80 > 192.168.122.1.58216: Flags [S.], seq
3306382458, ack 2182467642, win 65160, options [mss 1460,sackOK,TS val 2701277277
ecr 2517725285,nop,wscale 7], length 0
11:01:05.900035 IP 192.168.122.1.58216 > 192.168.122.123.80: Flags [.], ack 1, win 502,
options [nop,nop,TS val 2517725286 ecr 2701277277], length 0
```

2nd curl Hit

```
11:01:15.270105 IP 192.168.122.1.52764 > 192.168.122.123.80: Flags [S], seq 953647097,
win 64240, options [mss 1460,sackOK,TS val 2517734656 ecr 0,nop,wscale 7], length 0
11:01:15.270307 IP 192.168.122.123.52764 > 192.168.122.44.80: Flags [S], seq
953647097, win 64240, options [mss 1460,sackOK,TS val 2517734656 ecr 0,nop,wscale 7],
length 0
```

I have also done Load Testing:-

```
sudo apt install apache2-utils
```

Install ab command at local for load testing

install apache2-utils for ab command for Load Testing
run load test using ab command :-
ab -n 1000 -c 100 <http://192.168.122.123:80/>
jai@jai-Standard-PC-Q35-ICH9-2009:~\$ ab -n 1000 -c 100 http://192.168.122.123:80/
This is ApacheBench, Version 2.3 <\$Revision: 1843412 \$>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, <http://www.zeustech.net/>
Licensed to The Apache Software Foundation, <http://www.apache.org/>

Benchmarking 192.168.122.123 (be patient)

Completed 100 requests
Completed 200 requests
Completed 300 requests
Completed 400 requests
Completed 500 requests
Completed 600 requests
Completed 700 requests
Completed 800 requests
Completed 900 requests
Completed 1000 requests
Finished 1000 requests

Server Software: nginx/1.18.0
Server Hostname: 192.168.122.123
Server Port: 80

Document Path: /
Document Length: 234 bytes

Concurrency Level: 100
Time taken for tests: 0.184 seconds
Complete requests: 1000
Failed requests: 0
Total transferred: 475000 bytes
HTML transferred: 234000 bytes
Requests per second: 5434.22 [#/sec] (mean)
Time per request: 18.402 [ms] (mean)
Time per request: 0.184 [ms] (mean, across all concurrent requests)
Transfer rate: 2520.76 [Kbytes/sec] received

Connection Times (ms)

	min	mean	mean[+/-sd]	median	max
Connect:	1	8	4.1	7	28
Processing:	2	9	3.0	9	28
Waiting:	1	7	2.4	6	13
Total:	9	17	5.8	16	36

Percentage of the requests served within a certain time (ms)

50% 16
66% 17
75% 17
80% 18
90% 22
95% 33
98% 35
99% 35
100% 36 (longest request)

Output curl

curl http://192.168.122.123

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Hello Deepak!</title>
</head>
<body>
  <h1>Hello Deepak!</h1>
</body>
</html>
```

curl http://192.168.122.123

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Hello jaiprakash!</title>
</head>
<body>
  <h1>Hello jaiprakash!</h1>
</body>
</html>
```

